

# Scaling the L2:

EL mempool  
short term future

**Csaba Kiraly - EF Geth**

June 10th, 2025

# Scaling the L2:

EL mempool  
**present,**  
short term future

**Csaba Kiraly - EF Geth**

June 10th, 2025

# How blobs propagate in the mempool today?

## Mempool transaction diffusion 101

### a. Push

- i. Depends on **TX type** ( $\neq 3$ ) and **size** ( $< 4096$  bytes)
- ii. Pushing to **sqrt**(peercount) peers selected randomly
- iii. “**Stable**” **random** peer selection, based on **<tx signer address, self ID, and peer ID>**

### b. Notify and Pull

- i. Notify **all** (other) peers about tx: **<txhash, type, size>**
- ii. Pull **scheduling** policy

# How blobs propagate in the mempool today?

Measuring blob diffusion from a single Geth node


- For any given TX:

- a. Did we get it?
  - i. Push or pull, timestamped
- b. Did we see our peers having it?
  - i. Their announcement, receive timestamped
  - ii. From 50 peers .... from 500 peers
- c. Did we forward it?


- When to evaluate?

Focus on “Block transactions”, evaluate when we get the block with the TX

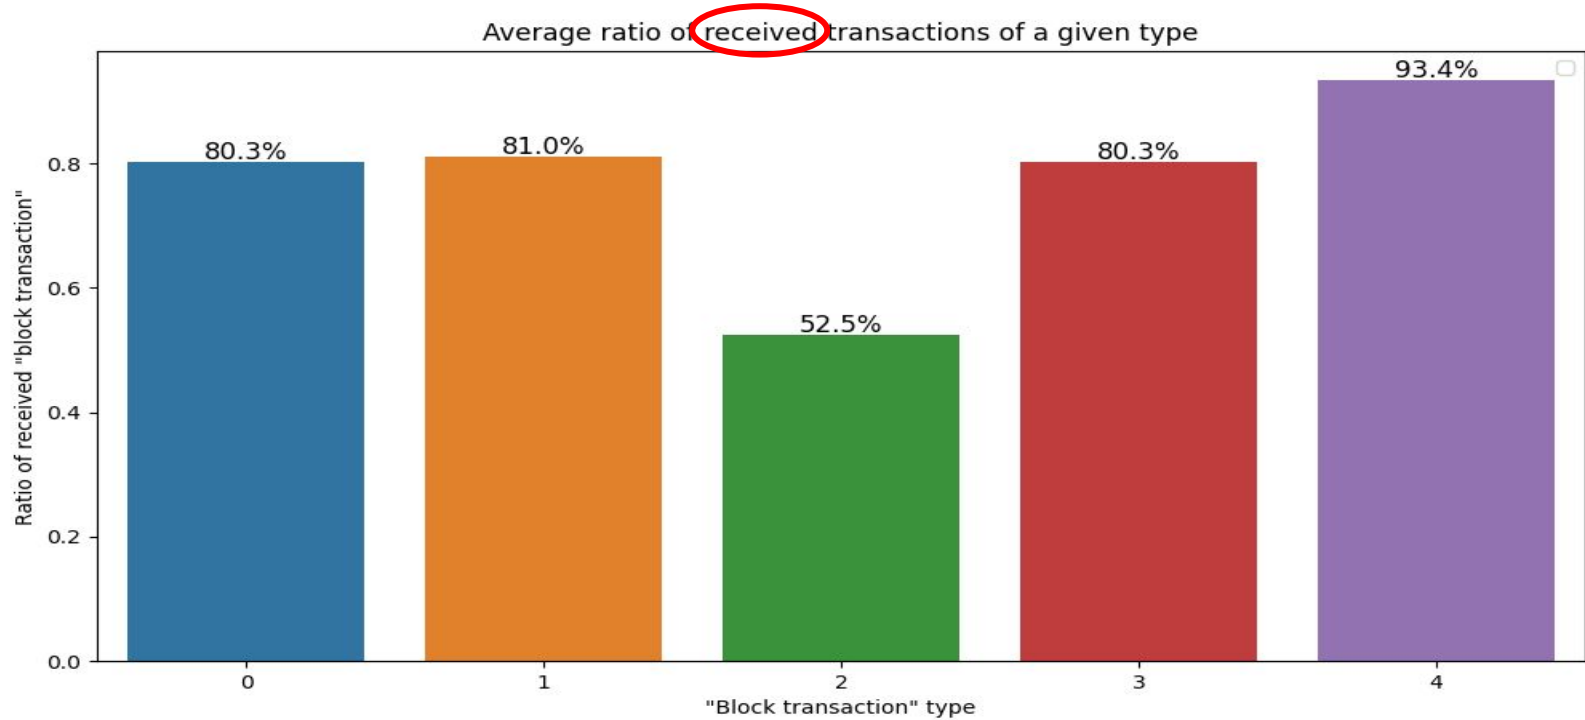
Overall ratio of transactions per type



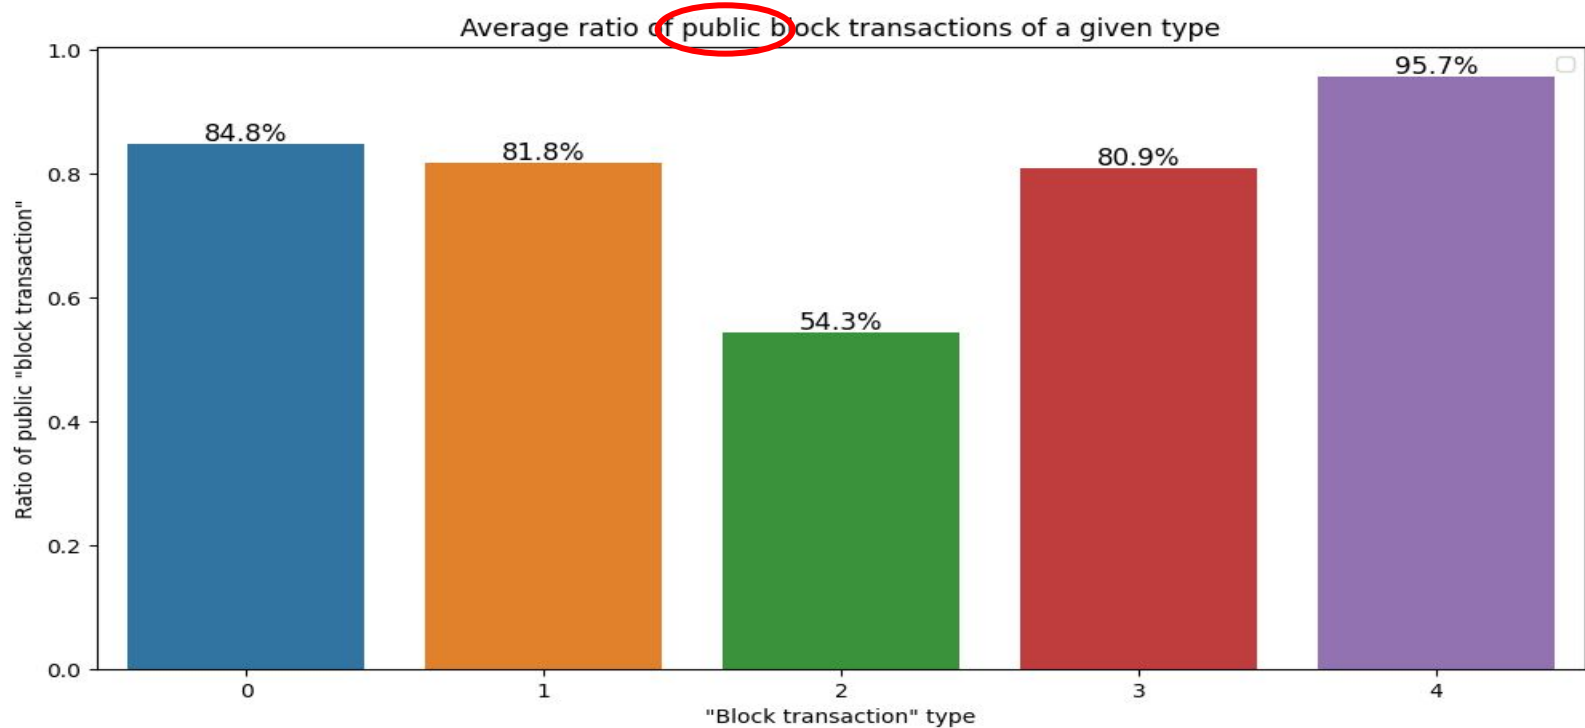
Overall ratio of transactions per type, weighted by transaction size



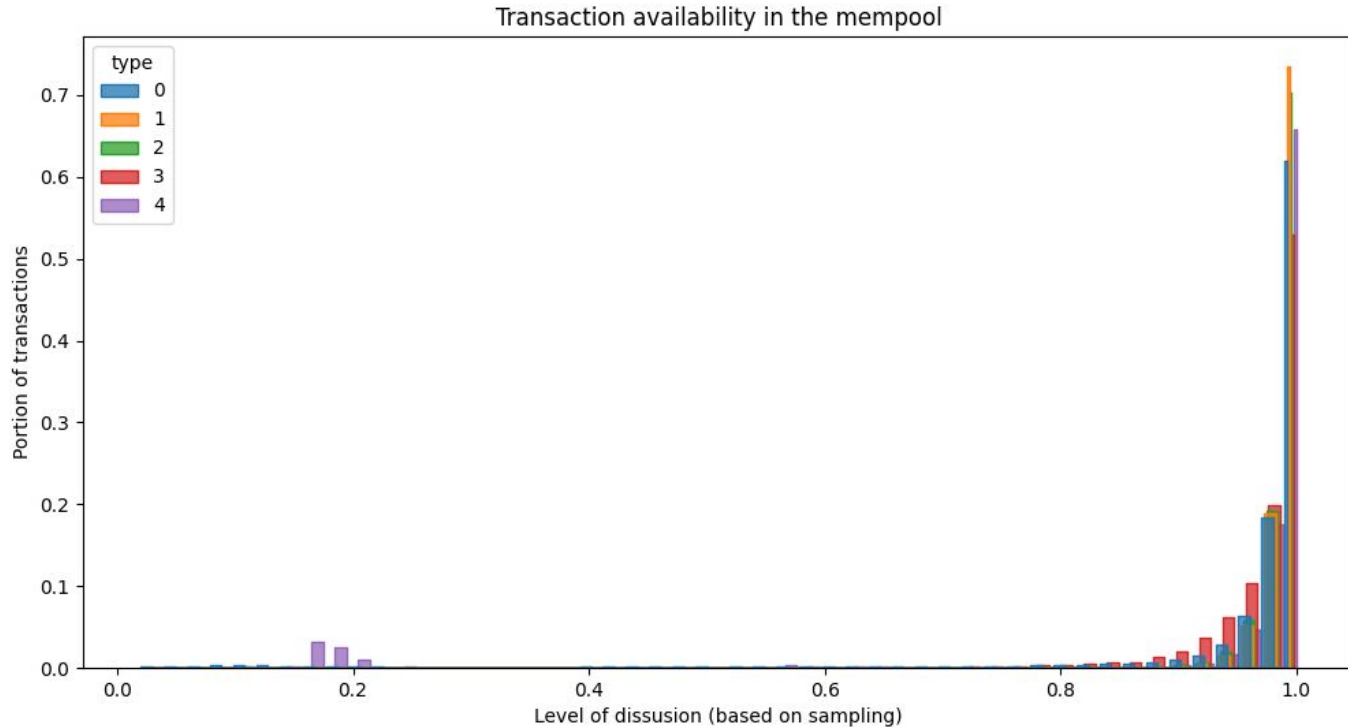
# Received vs. Public



# Received vs. Public

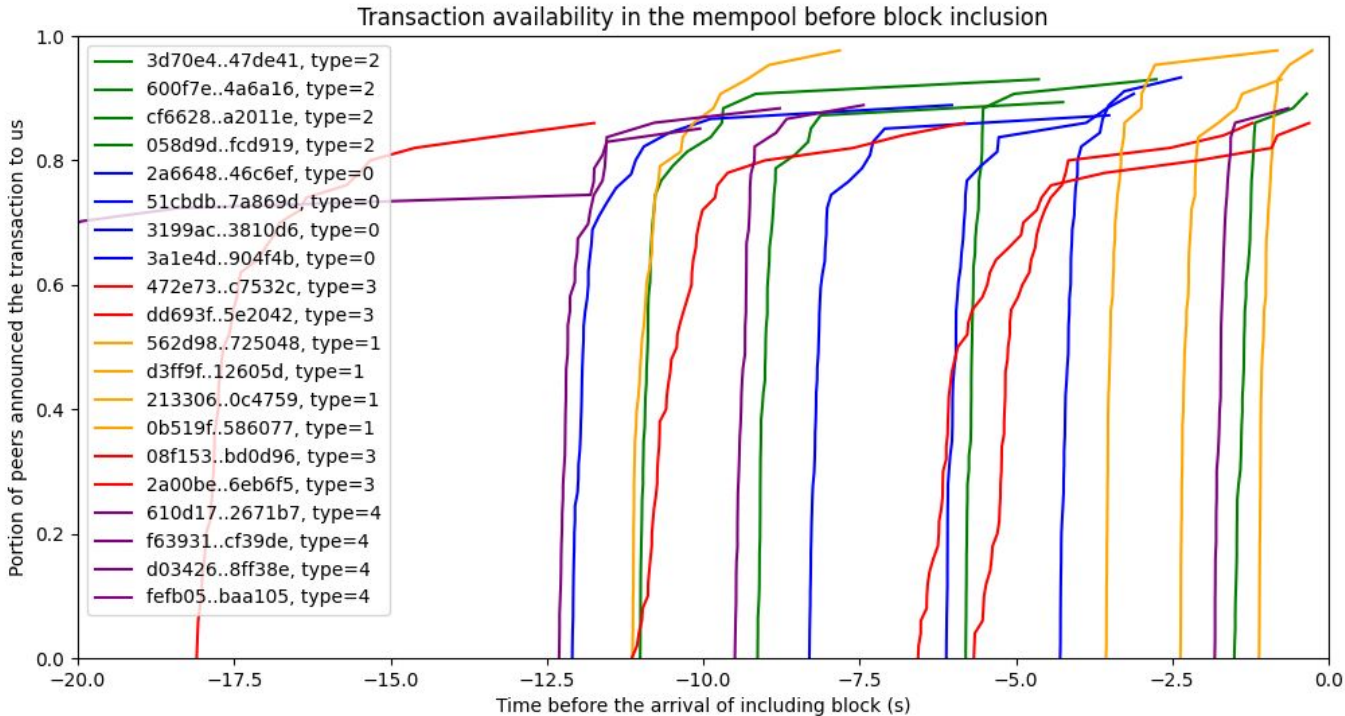


# Diffusion ratio (0..1) of transactions in the mempool

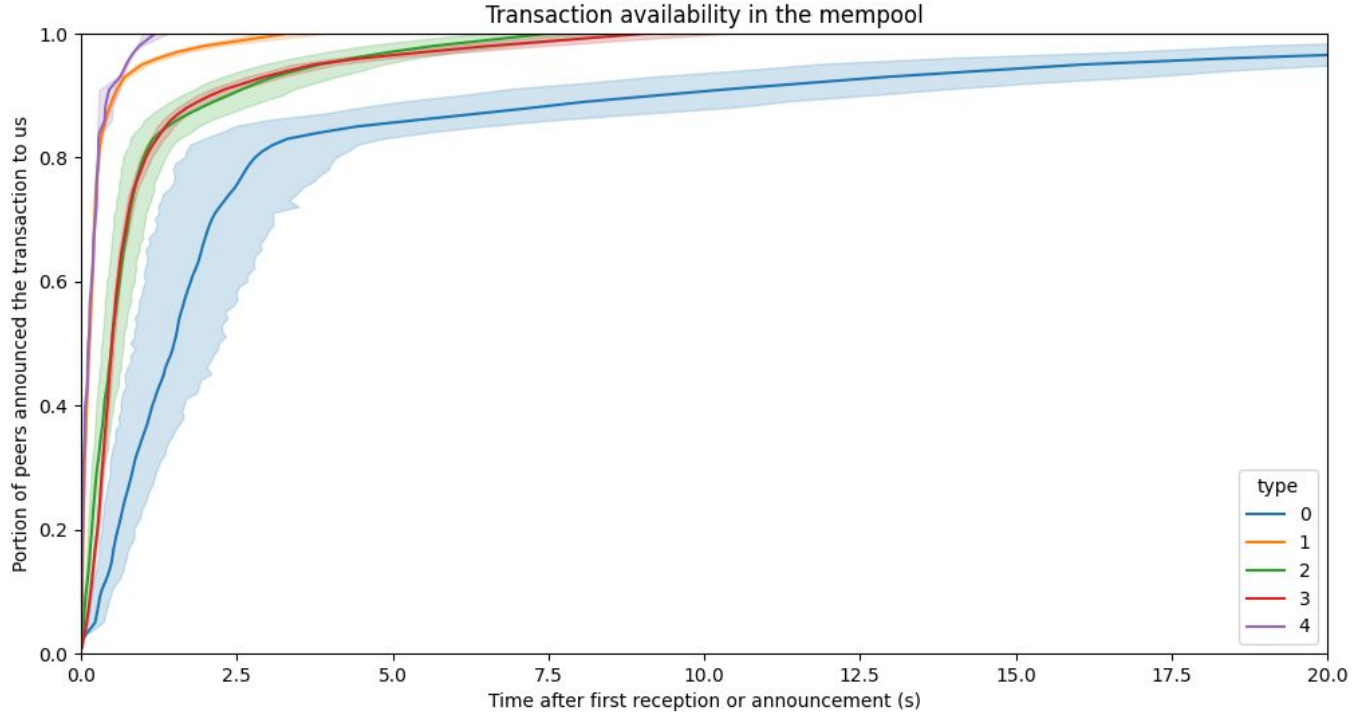




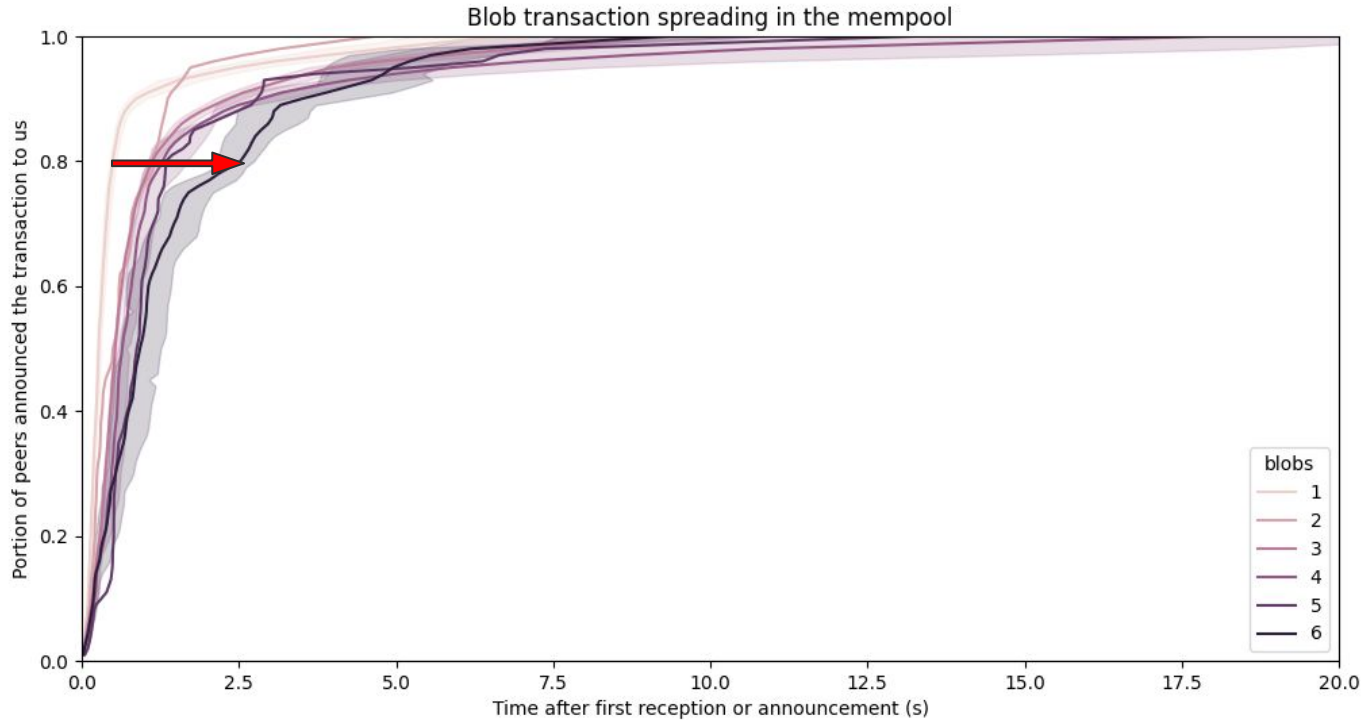
# Transaction “diffusion speed” in the mempool



# Transaction “diffusion speed” stats



# Transaction “diffusion speed” stats

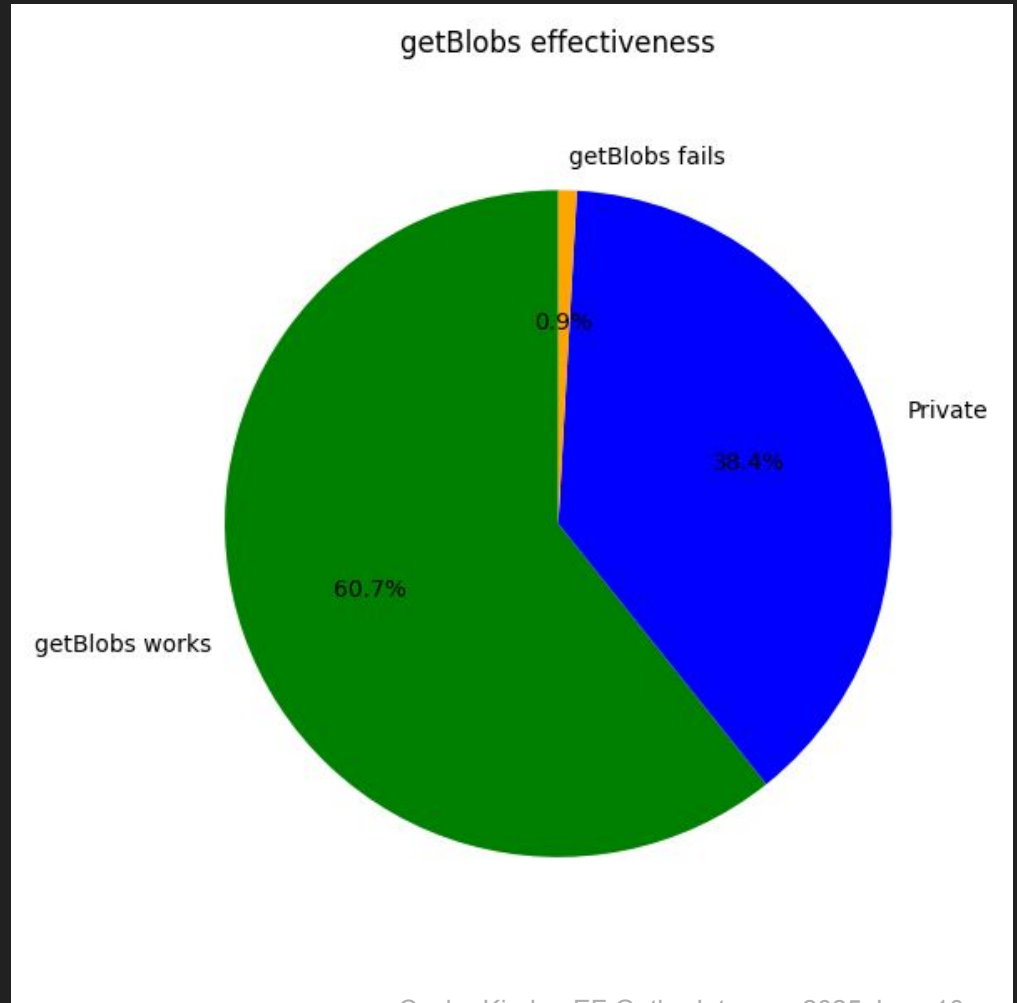


# GetBlobs effectiveness

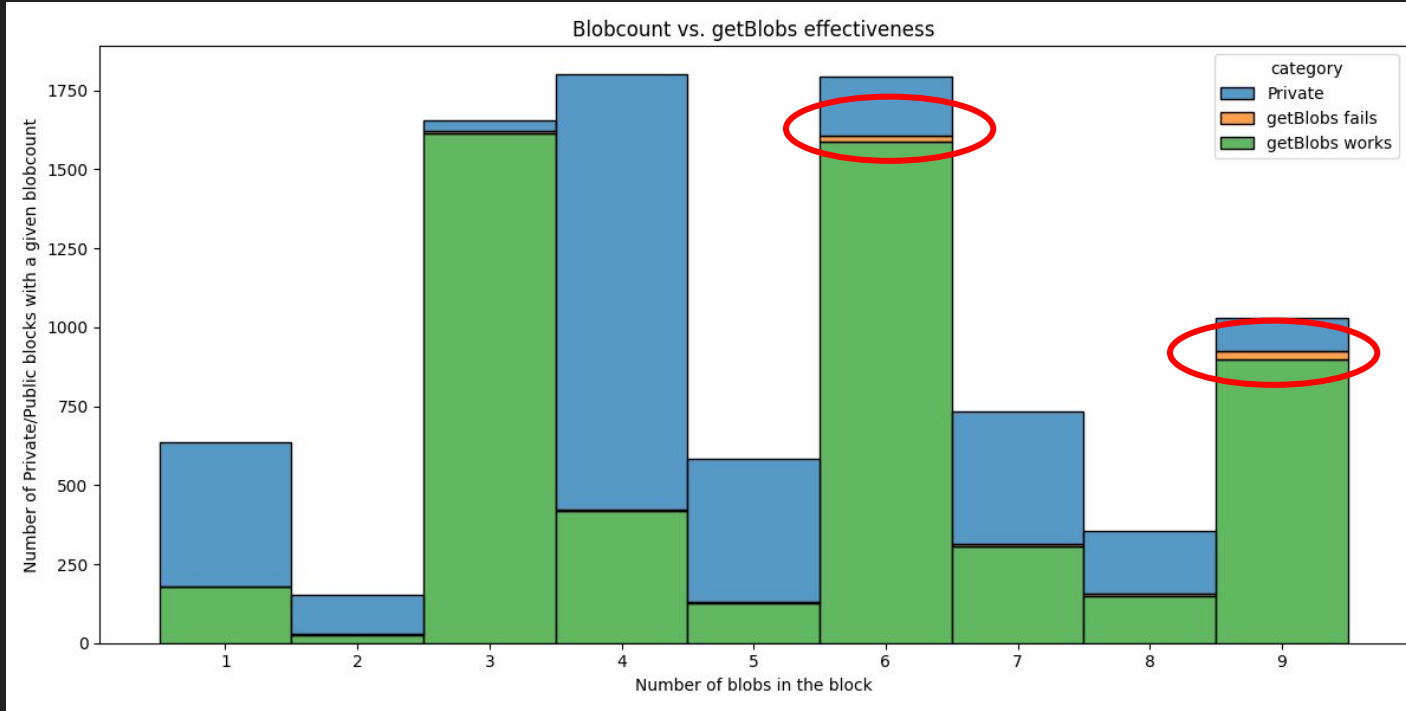
**Private:** it simply can't work,  
**choice of builder**

**Works:** our node had all the blobs  
by the time of block arrival

**Fails:** all blobs are public, but we  
miss some



# How would this scale with blob count?



Not enough info  
here

But it feels OK

**11 KB/sec/blob**

# Scaling the L2:

EL mempool  
present,  
**short term future**

**Csaba Kiraly - EF Geth**

**June 10th, 2025**

# How blobs should propagate in the mempool **tomorrow?**

Fully synched vs. "sharded" (implicitly or explicitly)

## **Today**

- the mempool feels basically **fully synched**

## **Tomorrow**

- If we scale **WITHOUT changing** anything, it will be **implicitly sharded** (segmented)
  - Current policy is **TX sender address based sharding** (not yet applied to blobs)
  - We have **no practical experience** yet of how this would play out
  - We just know **it will happen** as an interaction of resource limits, announcement and pull scheduling policies

# How blobs should propagate in the mempool **tomorrow**?

We can also introduce

- **Explicit horizontal sharding**, based on
  - TX hash
  - TX sender address
- **Vertical sharding**, based on blob segmentation and encoding
- **Sampling** in the mempool

Public mempool advantage (or what you miss with private blob feeds)

- pre-diffusion in the mempool (goes to DA and DAS through getBlobs today)
- sampling in the mempool?
- CR